


Inf264-Assignment 4

September 20, 2019

```
In [3]: import numpy as np
import matplotlib.pyplot as plt
from scipy import stats
from math import sqrt
from mpl_toolkits.mplot3d import Axes3D

# Some code (such as the code to do PCA) will generate useless warnings
# We will suppress these warning using the code below
import warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)

In [4]: from sklearn.datasets.samples_generator import make_blobs
Xy1 = np.genfromtxt('svmlin.csv', delimiter=',')
X1 = Xy1[:, :2]
y1 = Xy1[:, 2]
# Each datapoint in X is tuple of values (x0, x1). Lets plot them
# on x and y axes respectively.
plt.scatter(X1[:, 0], X1[:, 1], c=y1, s=50, cmap='autumn');
plt.xlabel(r"$x_0$")
plt.ylabel(r"$x_1$")
plt.grid()
```




0.1 Marginal distance from a hyperplane

Find two lines that classify the red X as two different classes without changin the label of any of the points.

```
In [5]: xfit = np.linspace(0, 3)
plt.scatter(X1[:, 0], X1[:, 1], c=y1, s=50, cmap='autumn')
plt.plot([0.6], [2.1], 'x', color='red', markeredgewidth=2, markersize=10)

# TODO
# plot two lines that classify X differently

for m, b in [(0.4, 2.2), (1.5, .4)]:
    plt.plot(xfit, m * xfit + b, '-k')
```



Find margin for both of you lines

```
In [6]: def calculate_margin(m,b,X):
        # TODO
        # find the minimum margin distance to the given line
        margins = []
        for i in X:
            margins.append(np.absolute(i[1]+m*i[0]+b)/(np.sqrt((m**2)+1)))# calculate dist
        return np.min(margins)

In [7]: # TODO
        # check the values of minimum margin for your lines
        print('margin of the first line: %.3f'%calculate_margin(.4, 2.2,X1))
        print('margin of the second line: %.3f'%calculate_margin(1.5, .4,X1))
```

```
margin of the first line: 2.669
margin of the second line: 1.452
```

Fit a support vector model to the dataset:

```
In [8]: from sklearn.svm import SVC # "Support vector classifier"
        model = SVC(kernel='linear', C=1E10)
        model.fit(X1, y1)

Out[8]: SVC(C=10000000000.0, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)
```


Complete the function below for plotting the SVM boundry:

```
In [9]: def plot_svc_decision_function(X, y, model, ax=None, plot_support=True):
        """Plot the decision function for a 2D SVC"""
        if ax is None:
            ax = plt.gca()
        xlim = ax.get_xlim()
        ylim = ax.get_ylim()

        # create grid to evaluate model
        x = np.linspace(xlim[0], xlim[1], 30)
        y = np.linspace(ylim[0], ylim[1], 30)
        Y, X = np.meshgrid(y, x)
        xy = np.vstack([X.ravel(), Y.ravel()]).T
        P = model.decision_function(xy).reshape(X.shape)
        # plot decision boundary and margins
        ax.contour(X, Y, P, colors='k',
                   levels=[-1, 0, 1], alpha=.5,
                   linestyles=['--', '-', '--'])

        # plot a black dot over the support vectors
        if plot_support:
            ax.scatter(model.support_vectors_[:, 0],
                      model.support_vectors_[:, 1],
                      s = 20, facecolor= 'black');
        ax.set_xlim(xlim)
        ax.set_ylim(ylim)

In [10]: plt.scatter(X1[:, 0], X1[:, 1], c=y1, s=50, cmap='autumn')
        plot_svc_decision_function(X1, y1, model, plot_support=True);
```




0.2 Kernel SVM


```
In [11]: Xy2 = np.genfromtxt('svmnonlin.csv',delimiter=',')
         X2 = Xy2[:, :2]
         y2 = Xy2[:, 2]

         plt.scatter(X2[:, 0], X2[:, 1], c=y2, s=50, cmap='autumn')

Out[11]: <matplotlib.collections.PathCollection at 0x2499af0ec50>
```



```
In [12]: clf = SVC(kernel='linear', C=1) # fit a linear SVM with C=1
clf.fit(X2, y2)
plt.scatter(X2[:, 0], X2[:, 1], c=y2, s=50, cmap='autumn')
plot_svc_decision_function(X2, y2, clf, plot_support=True);
```



```

In [13]: # Computes the weighted sum of kernel values
def compute_f(training_points , test_point, kernel, weights, gamma=1, intercept=0):
    f = 0
    for i in range(len(training_points)):
        trainSample = [training_points[i]] # put the training sample into two diemnsi
        f += weights[i]*kernel(test_point,trainSample,gamma) # Complete the dual form
    return f + intercept

In [14]: from sklearn.metrics.pairwise import polynomial_kernel, rbf_kernel
def surface_kernel(X,y,weights,clf,kernel):
    p = np.linspace(0, 6, 10)
    xx1, xx2 = np.meshgrid(p, p)

    z = float('NaN')*np.ones(xx1.shape)

    intercept = clf.intercept_
    for i in range(z.shape[0]):
        for j in range(z.shape[1]):
            tst = [[xx1[i, j], xx2[i, j]]]
            z[i, j] = compute_f(X, tst, kernel, weights, gamma=1, intercept=intercept)
    zz = np.empty(y.shape)

    for i in range(X.shape[0]):
        zz[i] = compute_f(X, [X[i,:]], kernel, weights, gamma=1, intercept=intercept)

    return xx1,xx2,z,zz

In [15]: def twodprojection(X,model):

    x_min = np.min(X[:, 0])
    y_min = np.min(X[:, 1])
    x_max = np.max(X[:, 0])
    y_max = np.max(X[:, 1])
    h = 0.01

    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                          np.arange(y_min, y_max, h))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])

    Z = Z.reshape(xx.shape)
    return xx,yy,Z

```

Fit a svm with RBF kernel to the dataset: (set the kernel value to 'rbf' and gamma to 1)


```

In [16]: clf = SVC(kernel='rbf', gamma=1) #Fit a SVM model with RBF kernel and gamma = 1
clf.fit(X2, y2)
weights = np.zeros((X2.shape[0], 1))
weights[clf.support_] = clf.dual_coef_.T

```

Plot the 3D projection of the surface of dual function of SVM:

```
In [17]: xx1, xx2, z, zz = surface_kernel(X2,y2,weights,clf,rbf_kernel)
fig = plt.figure(figsize=(10,8))
ax = fig.add_subplot(111, projection='3d')
ax.plot_surface(xx1, xx2, z, alpha=0.5)
ax.view_init(elev=45, azim=-30)
ax.scatter(X2[:, 0], X2[:, 1], zz, s=30,c=y2,cmap='autumn')
ax.set_xlabel('X1')
ax.set_ylabel('X2')
ax.set_zlabel('Z')
plt.show()
```




Plot the decision boundaries for RBF kernel:

```
In [18]: # Plot decision boundary
from matplotlib.colors import ListedColormap
cmap_light = ListedColormap(['#FFCCCC', '#80dfff'])
xx,yy,Z = twodprojection(X2,clf)
plt.figure()
plt.pcolormesh(xx, yy, Z, cmap=cmap_light)
```




```
plt.scatter(X2[:, 0], X2[:, 1], c=y2, s=50, cmap='autumn')
plt.show()
```



Fit a svm with RBF kernel to the dataset: (set the kernel value to 'poly' and gamma to 1 and degree to 2)

```
In [19]: clf_poly = SVC(kernel='poly', degree=2, gamma=1) # Fit a SVM with second degree polynomial kernel
clf_poly.fit(X2,y2)
weights = np.zeros((X2.shape[0], 1))
weights[clf_poly.support_] = clf_poly.dual_coef_.T
```


```
In [20]: xx1, xx2, z, zz = surface_kernel(X2,y2,weights,clf_poly,polynomial_kernel)
fig = plt.figure(figsize=(10,8))
ax = fig.add_subplot(111, projection='3d')
ax.plot_surface(xx1, xx2, z, alpha=0.5)
ax.scatter(X2[:, 0], X2[:, 1], zz, s=30, c=y2, cmap='autumn')
ax.set_xlabel('X1')
ax.set_ylabel('X2')
ax.set_zlabel('Z')
plt.show()
```



Plot the decision boundaries for Polynomial kernel:

```
In [21]: # Plot decision boundary
xx,yy,Z = twodprojection(X2,clf_poly)
plt.figure()
plt.pcolormesh(xx, yy, Z, cmap=cmap_light)
plt.scatter(X2[:, 0], X2[:, 1],c=y2, s=50, cmap='autumn')

plt.show()
```




0.2.1 Tuning the SVM: Softening Margins

Generate the dataset by running the code below:

```
In [22]: from sklearn.datasets.samples_generator import make_blobs

Xy3 = np.genfromtxt('svmmargin.csv', delimiter=',')
X3 = Xy3[:, :2]
y3 = Xy3[:, 2]
plt.scatter(X3[:, 0], X3[:, 1], c=y3, s=30, cmap='autumn');
plt.xlabel(r"$x_0$")
plt.ylabel(r"$x_1$")


Out[22]: Text(0,0.5,'$x_1$')
```



Fit two different SVMs with $C=10$ and $C=0.1$:

```
In [23]: fig, ax = plt.subplots(1, 2, figsize=(8, 8))
fig.subplots_adjust(left=0.0625, right=0.95, wspace=0.1)

for axi, C in zip(ax, [10.0, 0.1]):
    model = SVC(kernel='linear', C=C).fit(X3,y3) # fit a linear SVM with C=C
    axi.scatter(X3[:, 0], X3[:, 1], c=y3, s=50, cmap='autumn')
    plot_svc_decision_function(X3,y3,model, axi)
    axi.set_title('C = {0:.1f}'.format(C), size=14)
```



Load the iris dataset and split it into training and test sets and perform cross validation on the training set: ##### Hints: Use `train_test_split` and `cross_val_score` from `sklearn.model_selection` to split the dataset and perform cross validation for best value of `C`, respectively.

```
In [24]: from sklearn.model_selection import cross_val_score, train_test_split
         from sklearn import datasets
         iris = datasets.load_iris()
         # TODO
         # Split the dataset into training and test set. (75% to 25%)
         X_train, X_test, y_train, y_test = train_test_split( iris.data, iris.target, test_s
         score_list = []
         C_list = np.arange(0.1, 20, 0.05)
         # TODO
         # perform cross validation on the training set
         for C in C_list:
             clf = SVC(kernel='linear', C=C)
             score = cross_val_score(clf, X_train, y_train)
```

```

        score_list.append(score.mean())
    best_C = C_list[np.argmax(score_list)]
    print('The best value for C is :%.3f'%best_C)

```

The best value for C is :1.200

Measure the accuracy score on the test set

```

In [25]: from sklearn.metrics import accuracy_score
        clf = SVC(kernel='linear', C=best_C)
        clf.fit(X_train,y_train)
        print(accuracy_score(y_test,clf.predict(X_test)))

```

0.9736842105263158

0.3 Kernel Trick

Complete the implementation of phi:

```

In [26]: def phi(x):
        # TODO
        # Complete the implementation
        transformed_x = np.empty(0)
        for i in range(len(x)):
            transformed_x = np.concatenate((transformed_x,x[i]**2))# Power of two terms
        for i in range(len(x)):
            for j in range(len(x)):
                if (i!=j):
                    transformed_x = np.concatenate((transformed_x,np.sqrt(2)*x[i]*x[j]))
        for i in range(len(x)):
            transformed_x = np.concatenate((transformed_x,np.sqrt(2)*x[i])) #single Terms
        transformed_x = np.concatenate((transformed_x,[1]))
        return transformed_x

```

Measure the run time in normal way:

```

In [27]: import timeit
        np.random.seed(3)
        run_time = []
        kernel_values = []
        for d in range(1,500,100):
            start = timeit.default_timer()

            x = np.random.uniform(0,100,[d,1])
            z = np.random.uniform(0,100,[d,1])
            x_k = phi(x)
            z_k = phi(z)

```

```

kernel_values.append(np.transpose(x_k)*z_k) # the value of K(x,z) using phi

stop = timeit.default_timer()
run_time.append(stop - start)

```

Measure the run time for Kernel Trick:

```

In [34]: np.random.seed(3)
run_time_trick = []
kernel_values_trick = []
for d in range(1,500,100):
    start = timeit.default_timer()

    x = np.random.uniform(0,100,[d,1])
    z = np.random.uniform(0,100,[d,1])
    # TODO
    kernel_values_trick.append(((1+np.transpose(x)*z)**2)) # the value of K(x,z) using

    stop = timeit.default_timer()
    run_time_trick.append(stop - start)

```

Plot the run times:

```

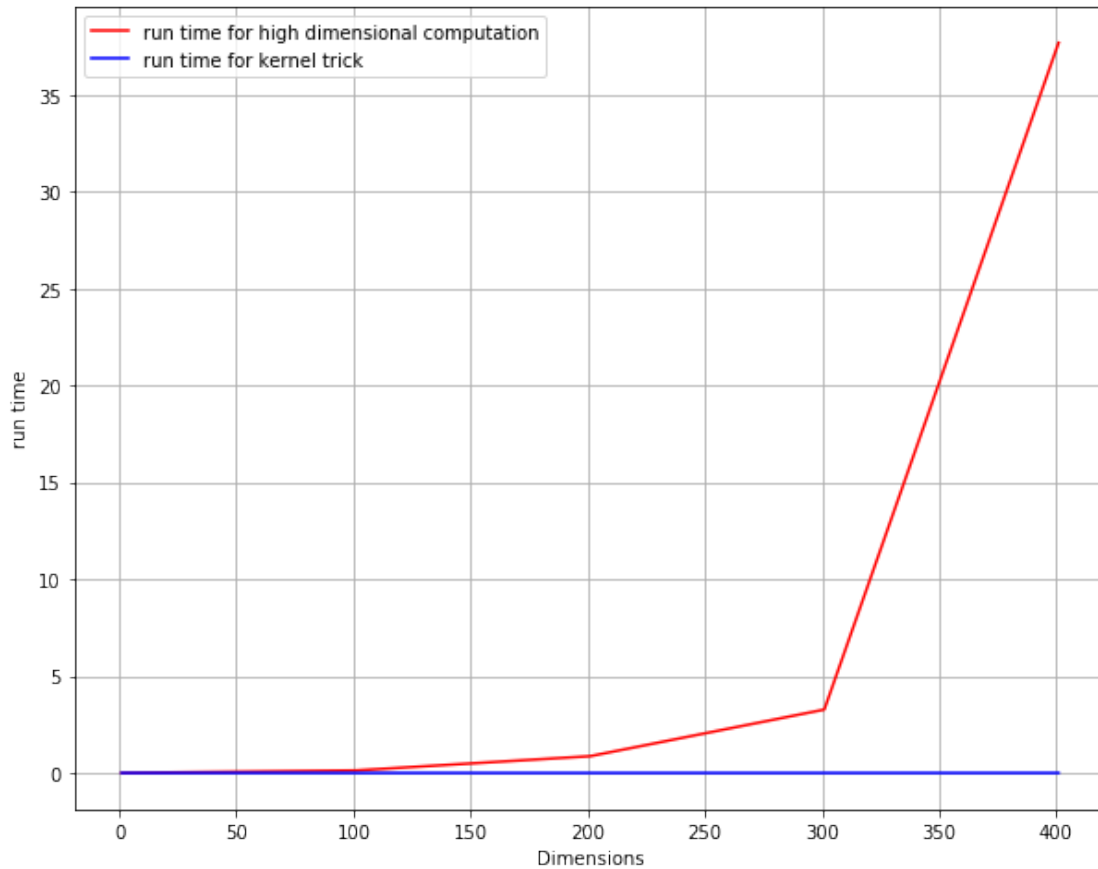
In [35]: plt.figure(figsize=(10,8))
plt.plot(range(1,500,100),run_time,'R',label='run time for high dimensional computation')
plt.plot(range(1,500,100),run_time_trick,'B',label='run time for kernel trick')
plt.grid()
plt.xlabel('Dimensions')
plt.ylabel('run time')
plt.legend()

```

```

Out[35]: <matplotlib.legend.Legend at 0x2499cadc208>

```



Compare the values of the kernel:

In [31]: kernel_values

```
Out [31]: [array([1.52136144e+07, 7.80092671e+03, 1.00000000e+00]),
          array([7.99690646e+06, 1.88963821e+07, 2.35614032e+07, ...,
                7.81970931e+03, 1.18629689e+03, 1.00000000e+00]),
          array([7.02100153e+05, 1.45283717e+07, 3.26126033e+06, ...,
                1.42011353e+04, 6.71088133e+03, 1.00000000e+00]),
          array([2.73471144e+07, 2.00567314e+07, 1.84968327e+06, ...,
                9.34249607e+03, 3.22712038e+02, 1.00000000e+00]),
          array([1.44058818e+07, 1.25365426e+06, 4.67468606e+06, ...,
                3.34306691e+03, 4.26562867e+03, 1.00000000e+00])]
```

In [32]: kernel_values_trick

```
Out [32]: [array([[3901.46335411]]),
          array([[2828.88020685, 4966.74679137, 8681.32273178, ..., 1868.13165033,
                7055.22254, 7631.35685001],
                [2476.51617897, 4347.99691059, 7599.72335084, ..., 1635.48034236,
                6176.24108182, 6680.58698807],
```



```

[1582.34165653, 2777.82988054, 4855.00898334, ..., 1045.09410616,
 3945.69890553, 4267.87138722],
...,
[ 223.53131035, 391.76412702, 684.07122313, ..., 147.9281661 ,
 556.1102842 , 601.44739666],
[1568.37337505, 2753.30163186, 4812.13262989, ..., 1035.87143103,
 3910.85465511, 4230.18132806],
[ 220.82625173, 387.01405447, 675.76790754, ..., 146.14212844,
 549.36244339, 594.14844442]]),
array([[ 838.91416791, 1990.58029721, 2270.65242188, ..., 3399.51294593,
 3329.33229323, 1934.76633396],
[1606.26424837, 3812.61011796, 4349.16837884, ..., 6511.82358872,
 6377.37247546, 3705.68250748],
[ 667.70377478, 1584.05080064, 1806.89599157, ..., 2705.09726492,
 2649.25657401, 1539.64126395],
...,
[ 484.67022861, 1149.44788888, 1311.11426684, ..., 1962.72781945,
 1922.21735474, 1117.23032603],
[1788.58179749, 4245.5129348 , 4843.01068722, ..., 7251.28900734,
 7101.56763719, 4126.44104347],
[1454.93537186, 3453.28816992, 3939.26487771, ..., 5898.04575131,
 5776.26939607, 3356.4406663 ]]),
array([[5230.44685194, 4189.10250652, 2035.8269159 , ..., 5832.04543233,
 5840.6963634 , 4702.65919317],
[5593.01760575, 4479.47422762, 2176.90660361, ..., 6236.32653466,
 6245.57725663, 5028.63709601],
[3496.23969912, 2800.22954745, 1361.03061483, ..., 3898.33408896,
 3904.11616851, 3143.47879919],
...,
[7022.16267105, 5624.03238063, 2733.00037947, ..., 7829.88127211,
 7841.49618971, 6313.54412822],
[4184.10162106, 3351.11692062, 1628.68415312, ..., 4665.32804292,
 4672.24803366, 3761.91750226],
[ 180.46913862, 144.7312914 , 70.83312847, ..., 201.11537179,
 201.41226268, 162.35601898]]),
array([[3796.50810202, 1444.10741239, 2984.89164979, ..., 957.23165485,
 1675.57334411, 2373.99119918],
[2945.82950752, 1120.66703174, 2316.11880395, ..., 742.91362989,
 1300.25503086, 1842.13808128],
[2751.19253487, 1046.66322239, 2163.10223231, ..., 693.87723491,
 1214.38144624, 1720.44901865],
...,
[4589.76558654, 1745.71545144, 3608.52208887, ..., 1157.08313633,
 2025.55753671, 2869.94404103],
[3789.61803244, 1441.48770767, 2979.47492545, ..., 955.49578644,
 1672.53345418, 2369.68345591],
[3412.35445283, 1298.04660479, 2682.88389872, ..., 860.44886061,
 1506.08524299, 2133.81433639]]])

```